

CHRONUS: A Strategic Research and Architectural Roadmap for Next-Generation Memory-Grounded Intelligence

1. Vision Summary

The irreversible loss of human cognitive identity at the end of life represents a profound information deficit in human history. Traditionally, personal archives—comprising diaries, letters, voice recordings, and digital communications—have remained static, searchable only by keyword and entirely devoid of the semantic interconnectedness that characterizes human memory. The baseline CHRONUS system was established to address this precise gap by translating fragmented personal archives into an interactive, memory-grounded, voice-authentic conversational agent.¹ By leveraging local vector databases, semantic sentence embeddings, and neural text-to-speech synthesis, the foundational architecture successfully demonstrated that an individual's documented thoughts and beliefs could be queried and re-animated without the hallucinatory risks inherent to unconstrained large language models.¹

However, to transition CHRONUS from a localized, proof-of-concept memory prototype into a research-grade, future-ready intelligence platform, the underlying architecture must fundamentally evolve. The next-generation CHRONUS vision transforms the system into a multi-paradigm, highly secure platform. It anchors itself firmly in the clinical, psychological, and ethical rigor of Grief Technology, ensuring that the simulation of deceased individuals prioritizes emotional safety and data sovereignty.² Concurrently, it scales securely via Privacy-Preserving Cloud MLOps, utilizing advanced cryptographic techniques like Additively Homomorphic Encryption to enable distributed memory retrieval without exposing plaintext data.⁴

Crucially, this evolution establishes dedicated research sandboxes for Quantum Artificial Intelligence (QML) and Neuromorphic (Spiking Neural Network) memory architectures. The objective is not to deploy speculative quantum or neuromorphic hardware into immediate production, as such claims would undermine scientific credibility. Rather, the goal is to architect a layered system where the classical retrieval-augmented pipeline operates robustly in the present, while hybrid quantum-classical algorithms and biologically plausible memory consolidation mechanisms can be tested, benchmarked, and eventually integrated as the underlying hardware matures.⁶ Ultimately, CHRONUS becomes a tiered reality model: a viable product today, an advanced cloud-scale platform tomorrow, and a pioneering research frontier for the next decade of cognitive simulation.

2. Three-Pillar Strategic Framing

The strategic evolution of CHRONUS is governed by a rigorous multi-pillar research framework. While the system operates across four distinct technical domains, the algorithmic and domain-specific advancements are structured around three primary research pillars, all supported by a foundational infrastructure layer. This framing isolates immediate product requirements from long-term scientific moonshots, ensuring the platform remains scientifically credible without succumbing to technological hype.

The first pillar is Grief Tech and Digital Legacy, which serves as the application and ethical layer. This pillar focuses on the clinical, psychological, and ethical deployment of simulated identities. It prioritizes factual provenance, the mitigation of maladaptive psychological dependencies, and strict adherence to deterministic memory retrieval over generative hallucination.² The second pillar, Quantum AI and Quantum Compute, represents the optimization frontier. Rather than claiming immediate quantum supremacy for end-user applications, this pillar establishes a research pathway investigating the transposition of classical text embeddings into quantum states, exploring amplitude encoding and hybrid parameterized quantum circuits (PQCs) to potentially accelerate high-dimensional clustering and similarity search in the future.⁷ The third pillar, Neuromorphic Computing and Neuroscience, constitutes the biological frontier. This pillar examines event-driven processing and spiking neural networks (SNNs) to model memory consolidation (semantization) through simulated sleep and awake phases, drawing direct inspiration from hippocampal-cortical interactions and Spike-Timing-Dependent Plasticity (STDP).⁶

Supporting these three algorithmic and domain pillars is the foundational infrastructure layer: Cloud-First Storage and Distributed MLOps. This layer addresses the urgent need to transition the local-only ChromaDB implementation into a distributed, zero-trust cloud architecture. It leverages advanced privacy-preserving methodologies, including Federated Learning, Differential Privacy, and Partially Homomorphic Encryption (PHE), to allow scalable memory storage and family collaboration without compromising the extreme privacy requirements of posthumous personal data.⁴

3. Literature Review by Pillar

3.1 Grief Tech and Digital Legacy Literature

The intersection of artificial intelligence and bereavement has generated significant clinical, ethical, and sociological debate. "Griefbots" or digital memorial agents utilize natural language processing to simulate a deceased individual's conversational patterns, drawing from digital footprints such as texts, emails, and voice recordings.³ Psychologically, grieving involves a complex process of reconciling the emotional presence of a loved one with the reality of their physical absence. Systems that produce overly convincing illusions of independent consciousness risk transitioning healthy "continuing bonds"—a process of adaptive meaning reconstruction following bereavement—into maladaptive dependencies.¹¹ Clinical literature warns that excessive reliance on simulated entities can foster conditions akin to Prolonged

Grief Disorder (PGD), characterized by intense, persistent sorrow and an inability to accept the finality of death, thereby isolating the user from authentic human support networks.² Ethicists and clinicians emphasize that the commercialization of digital afterlives introduces severe vulnerabilities surrounding posthumous autonomy and data commodification.² The taxonomy of "Digital Ghosts" proposed by Spitale and Germani (2025) outlines multidimensional ethical profiles based on the timing of creation (pre-mortem intentional legacy planning vs. post-mortem third-party creation), data sources, and the primary purpose of the avatar.¹³ A critical risk identified in the literature is "platform temporality"—the paradoxical situation wherein venture-funded startups promise infinite digital immortality but operate on volatile business models, risking a traumatic secondary loss of the digital entity if the platform ceases operations.² Consequently, clinical guidelines mandate strict UX disclosures framing these systems as deterministic memory retrieval tools rather than sentient entities, alongside privacy-preserving architectures that keep the digital remains exclusively under the control of the authorized estate.²

3.2 Quantum AI and QNLP Literature

Quantum Natural Language Processing (QNLP) seeks to exploit the principles of quantum mechanics to overcome classical computational bottlenecks in embedding and processing linguistic structures.¹⁵ Recent literature, such as the work by Cavar et al. (2025), investigates the mapping of classical word embeddings (e.g., Sentence-BERT) into gate-based quantum systems, exploring how semantic spaces can be transposed into quantum states.⁸

A primary technique explored in this domain is amplitude encoding, which maps an N -dimensional classical vector into the amplitudes of a $\log_2(N)$ -qubit quantum state.⁷ While this method is highly qubit-efficient and theoretically preserves inner-product geometry—allowing cosine similarity to be computed via inner products in Hilbert space—state preparation is computationally expensive, typically requiring $O(N)$ operations.⁷ Furthermore, the resulting states are highly sensitive to noise, making deep amplitude encoding circuits challenging to deploy on current Noisy Intermediate-Scale Quantum (NISQ) devices without robust error correction.⁷ To mitigate these limitations, current research focuses on hybrid approaches combining classical embeddings with Parameterized Quantum Circuits (PQCs) or quantum feature maps to exploit the high-dimensional Hilbert space structure for tasks like text classification and sentiment prediction.⁷ However, practical superiority over classical cosine similarity for standard semantic text retrieval remains largely speculative, positioning QNLP as a mid-to-long-term research vector rather than an immediate production enhancement.⁷

3.3 Neuromorphic Computing and Neuroscience Literature

Neuromorphic computing addresses the significant latency and energy costs associated with traditional deep neural networks (DNNs) by employing biologically inspired Spiking Neural Networks (SNNs).⁶ These architectures utilize simplified, highly efficient approximations of

biological neurons, such as the Leaky Integrate-and-Fire (LIF) model, to encode information through sparse spike events and temporal dynamics, naturally supporting temporal coding and event-driven processing.⁶

Recent developments in computational neuroscience involve bridging complex biological memory models to specialized neuromorphic hardware, such as Intel's Loihi 2 chip, utilizing software frameworks like Lava.⁶ A prime example is the Hippocampal-Cortical (HiCo) SNN, which simulates the interaction between hippocampal, perceptual, and neocortical areas to model memory semantization—the process by which episodic memories are consolidated into generalized semantic knowledge.⁶ The learning dynamics of the HiCo model follow a two-stage process inspired by biological memory consolidation, reproducing both awake encoding phases and offline sleep consolidation phases.⁶ This is achieved through a three-factor Spike-Timing-Dependent Plasticity (STDP) rule driven by a neuromodulatory signal (analogous to dopamine), which reinforces causal relationships between neural activity patterns.⁶ This biologically grounded approach provides a stark contrast to the static vector spaces of classical transformers, offering a pathway toward memory systems that continuously learn, associate, and adapt over time.

3.4 Cloud MLOps and Privacy-Preserving Architecture Literature

Transitioning sensitive personal memory archives from localized, air-gapped environments to scalable cloud infrastructures introduces severe data leakage risks. Exposed vector databases have historically resulted in the catastrophic compromise of Personally Identifiable Information (PII), corporate credentials, and highly sensitive internal communications.²⁰ To mitigate these vulnerabilities in distributed machine learning systems, literature on privacy-preserving cloud architectures, such as the PACC-Health framework, advocates for a tightly coupled, multi-layered approach.⁹ This involves integrating Federated Learning (FL) for decentralized model training, Differential Privacy (DP) to mathematically bound privacy leakage by injecting calibrated noise into gradients, and Zero-Knowledge Proofs (zk-SNARKs) to cryptographically verify policy compliance without exposing the underlying plaintext data.²¹

Furthermore, to perform vector similarity searches securely in the cloud, research highlights the critical role of Homomorphic Encryption (HE). While Fully Homomorphic Encryption (FHE) allows arbitrary computations on encrypted data, it introduces massive computational overhead—specifically through bootstrapping and ciphertext-ciphertext multiplication—making it impractical for large-scale, real-time vector retrieval.⁵ Recent cryptographic advancements demonstrate that Additively Homomorphic Encryption (AHE) or Partially Homomorphic Encryption (PHE) can efficiently compute encrypted similarities.⁴ Because AHE only supports addition and scalar multiplication, direct cosine similarity calculations are mathematically impossible; however, by pre-normalizing both the source and query vectors locally, the cosine similarity is reduced to a simple dot product, which can be executed homomorphically with significantly reduced ciphertext sizes and latency.⁴ Finally, to manage the dynamic trade-offs between privacy, accuracy, and latency, modern MLOps architectures utilize Reinforcement Learning (RL) controllers to adaptively govern differential privacy budgets and access policies in real time.⁹

4. Feasibility Matrix

The following matrix categorizes proposed next-generation features by their current scientific and engineering reality, establishing clear boundaries between immediate product enhancements, mid-term research prototypes, and speculative moonshots.

Feature / Architectural Component	Feasibility Classification	Technical Rationale
Semantic memory retrieval	Feasible Now	Core baseline functionality effectively operating using Sentence-BERT and local ChromaDB. ¹
Voice cloning (Local & Cloud)	Feasible Now	Fully implemented via XTTS-v2 for privacy-preserving local generation and ElevenLabs API for high-fidelity cloud generation. ¹
Multimodal ingestion	Feasible Now	Standard OCR (pdfplumber) and audio-to-text transcription pipelines are readily available and integrated into standard Python workflows. ¹
Grief safety layer & provenance	Feasible Now	Achievable through strict UX design, rule-based clinical disclaimers, and explicit source-citation overlays injected into the frontend. ²
Cloud vector memory (PHE/AHE)	Feasible Prototype	Additively Homomorphic Encryption for secure cosine similarity is mathematically viable by pre-normalizing vectors, though it requires

		significant algorithmic refactoring of the database layer. ⁴
MLOps versioning & Data Lineage	Feasible Now	Standard industry practices using tools like MLflow or DVC can track model lineage and dataset versions effectively.
Federated memory collaboration	Feasible Prototype	Federated Learning allows family members to train models locally and share only secure aggregated updates, though orchestration across consumer edge devices is complex. ⁹
Avatar interaction	Feasible Prototype	Visual avatars (e.g., SadTalker) can be integrated, but real-time conversational latency over WebSockets remains an engineering challenge. ²
Grounded Response (Mix Method)	Feasible Now	Deterministic templating combined with cosine-similarity thresholds effectively eliminates generative ideological hallucination. ¹
Confidence scoring	Feasible Now	Distance thresholds (e.g., $\theta = 1.45$) generated by HNSW indexing map directly to explainable confidence percentages. ¹
Quantum-inspired retrieval	Feasible Prototype	Tensor networks running on classical GPUs can simulate

		quantum clustering and optimization algorithms to improve ranking speed.
Brain-inspired consolidation (HiCo)	Feasible Prototype	Feasible offline on Intel Loihi 2 via the Lava framework for experimental memory semantization during simulated "sleep" cycles. ⁶
Quantum retrieval (Gate-based)	Speculative Future	Deep amplitude encoding circuits require fault-tolerant, post-NISQ hardware to overcome noise and state preparation overhead. ⁷
Neuromorphic memory recall	Speculative Future	Integrating SNN temporal recall directly into real-time, interactive NLP query pipelines is currently unrealistic due to massive hardware and software abstraction gaps.

5. Gap Analysis

Current conversational AI and digital legacy systems fail on several critical fronts that the CHRONUS architecture uniquely resolves. By analyzing the market against the multi-pillar strategy, specific technological and ethical gaps become apparent.

The first major deficiency in existing platforms is the **Hallucination Gap**. Systems like Replika, or standard large language models operating under unconstrained persona prompts, generate text by predicting the most statistically probable continuation given a context window. This architecture fundamentally leads to factual fabrications, where the AI invents memories, opinions, and ideological stances that the deceased individual never held.¹ For a grief-sensitive system, this is not merely an inaccuracy; it is a violation of posthumous autonomy. CHRONUS resolves this through its "Mix Method"—a deterministic, retrieval-augmented prototype that anchors every generated response to an explicitly retrieved sentence from the vector database, entirely eliminating ideological hallucination.¹

The second major deficiency is the **Privacy and Commodification Gap**. Commercial entities frequently require the ingestion of highly intimate personal data into proprietary cloud environments. This raises severe ethical risks regarding data commodification and "platform

temporality"—the traumatic loss of a simulated entity if the host company goes bankrupt.² Existing systems lack cryptographic guarantees against data leakage.²⁰ CHRONUS currently solves this via a strict local-first execution baseline.¹ The next-generation architecture bridges the gap to cloud scalability by implementing Additively Homomorphic Encryption (AHE), allowing the system to scale storage to the cloud while keeping the decryption keys entirely on the user's edge device, thereby rendering the cloud data useless to corporate entities or attackers.⁴

The final deficiency is the **Hardware Evolution Gap**. Classical vector databases are inherently static; they index embeddings mathematically but do not natively consolidate, forget, or re-weight information over time based on associative usage.⁶ They ignore the biological realities of memory. Existing platforms treat AI as a pure software problem, ignoring hardware horizons. The introduction of the Neuromorphic (SNN) and Quantum (QML) pillars positions CHRONUS to transition from static indexing to dynamic, biologically plausible memory semantization.⁶ By establishing isolated research sandboxes today, CHRONUS prepares to swap its classical vector search for quantum-accelerated retrieval and neuromorphic associative recall as those nascent hardware platforms mature.⁷

6. Proposed Architecture

To support the multi-pillar research strategy while maintaining a robust, production-grade product, the CHRONUS architecture must be decoupled into distinct operational planes. This layered approach isolates experimental algorithms from user-facing systems.

6.1 The Application and Presentation Plane (Frontend)

The user interface is built on React 18 and TypeScript, focusing on emotional safety and clinical transparency.

- **Authentication and Dashboard:** Manages secure login, persona library browsing, and access to the custom model builder.¹
- **Ethics & Safety Gateway:** A middleware layer that filters outputs for emotional distress markers, enforces clinical session limits, and appends explicit source citations (provenance overlays) to every response, framing the system as an archive rather than a sentient being.²
- **Multimodal Ingestion Pipeline:** Handles drag-and-drop file uploads (PDF, EML, TXT, DOCX), interview augmentation, and structural validation.¹

6.2 The Classical Inference & Memory Plane (Product-Grade)

This layer represents the immediate, highly stable core of the CHRONUS product, executed locally or on trusted edge nodes.

- **Backend Orchestration:** Python/Flask APIs manage the asynchronous processing of document parsing and text extraction using libraries like spaCy and pdfplumber.¹
- **Semantic Memory Engine:** Sentence-BERT (all-MiniLM-L6-v2) generates 384-dimensional dense vectors.¹ For local execution, ChromaDB provides extremely fast

$O(\log n)$ retrieval using Hierarchical Navigable Small World (HNSW) indexing.¹

- **Response Synthesis:** The deterministic Mix Method constructs responses by wrapping retrieved authentic quotes with persona-specific thematic explanations.¹
- **Voice Engine:** XTTS-v2 operates completely locally to synthesize neural speech without transmitting biometric voice data over the network.¹

6.3 The Secure Cloud MLOps Plane (Mid-Term Scalability)

To support multi-user family collaboration and massive archival datasets, CHRONUS expands into a zero-trust cloud architecture.

- **Encrypted Vector Store:** Transitions from local ChromaDB to a cloud-hosted vector database secured by Additively Homomorphic Encryption (AHE). Sentence embeddings are normalized on the edge device; the cloud executes encrypted dot products without ever decrypting the data.⁴
- **Adaptive Governance (RL Controller):** Inspired by the PACC-Health framework, a Proximal Policy Optimization (PPO) reinforcement learning agent continuously monitors telemetry (latency, privacy leakage metrics) to dynamically adjust differential privacy noise levels and Zero-Knowledge Proof (ZKP) access policies across the distributed network.⁹
- **Observability & Versioning:** MLOps tools (e.g., MLflow) manage dataset versioning, model rollback, and secure federated aggregation of new memories.⁹

6.4 The Neuromorphic Research Sandbox (Speculative / Offline)

This isolated environment operates entirely offline for computational neuroscience research.

- **Hardware / Framework:** Intel Loihi 2 hardware accessed via the Lava software framework.⁶
- **Function:** Tests Hippocampal-Cortical (HiCo) SNN models. It ingests batches of classical SBERT embeddings, converts them into temporal spike trains using Leaky Integrate-and-Fire (LIF) neurons, and runs simulated offline "sleep" cycles.⁶ Using three-factor Spike-Timing-Dependent Plasticity (STDP), the sandbox observes associative memory consolidation, studying how a simulated brain might cluster related memories over time.⁶

6.5 The Quantum Research Sandbox (Speculative / Offline)

An environment dedicated to benchmarking Quantum Natural Language Processing (QNLP).

- **Hardware / Framework:** Qiskit or PennyLane running on classical simulators or available NISQ backends.
- **Function:** Experiments with amplitude encoding to map reduced-dimensionality classical vectors into n -qubit quantum states.⁷ The sandbox benchmarks Quantum Support Vector Machines (QSVMs) and Parameterized Quantum Circuits (PQCs) against classical cosine similarity to identify future advantages in high-dimensional semantic clustering.⁷

7. Data Flow

The CHRONUS data flow dictates how raw, unstructured artifacts transform into securely retrieved, voice-cloned intelligence, incorporating the necessary cryptographic and mathematical transformations required for the next-generation architecture.

1. **Ingestion & Preprocessing:** The user uploads legacy data through the React frontend. Local Python parsers (e.g., pdfplumber, email extractors) clean the text, removing HTML tags, metadata, and whitespace formatting.¹
2. **Chunking & Embedding:** spaCy segments the normalized text into discrete sentence units based on linguistic boundaries. Each unit is processed by the all-MiniLM-L6-v2 transformer, converting the text into a dense vector $v_i \in \mathbb{R}^{384}$.¹
3. **Normalization & Encryption (Cloud Path):** If the user opts into secure cloud storage, the system performs a min-max normalization on the vectors. Additively Homomorphic Encryption (AHE) is then applied locally using a public key, encrypting the vectors before transmission to the cloud.⁴
4. **Storage:** For local deployment, plaintext vectors and metadata are stored in ChromaDB.¹ For cloud deployment, the ciphertext vectors are stored in a distributed vector database (e.g., AHE-enabled Milvus or pgvector).²⁵
5. **Query Processing & Retrieval:** A user query is embedded into a vector q . In the local path, cosine similarity $\cos(q, d) = \frac{q \cdot d}{\|q\| \times \|d\|}$ is calculated directly via ChromaDB.¹ In the cloud path, the query vector is encrypted and sent to the server. The server calculates the dot product homomorphically on the ciphertexts, returning the encrypted distance scores. The local client decrypts the scores to identify the top- k matches.⁴ Results are filtered by a strict distance threshold ($\theta = 1.45$) to establish confidence scoring and reject irrelevant data.¹
6. **Response Synthesis & Safety Check:** The Mix Method constructs the response: [Persona Intro] + + +.¹ The Ethics Gateway scans the output for clinical safety flags.
7. **Delivery & Feedback:** The approved text is routed to XTTS-v2 for local acoustic synthesis. The frontend displays the text alongside explicit source citations (e.g., "From Diary, 1998") and plays the audio stream.¹ User feedback (e.g., marking a response as inaccurate) is logged locally for future dataset pruning or retraining.
8. **Offline Research Sync:** Periodically, non-sensitive vector clusters and metadata are exported to the isolated Neuromorphic and Quantum sandboxes for experimental semantization and amplitude encoding benchmarks.⁶

8. Product Roadmap

The evolution from a local prototype to a multi-paradigm research platform follows a rigorous five-phase roadmap.

- **Phase 1: Present CHRONUS Hardening (Months 1-3)** The immediate priority is solidifying the local-first product baseline. Engineering efforts will focus on optimizing the ChromaDB and XTTS-v2 pipelines to reduce inference latency. Crucially, this phase implements the primary Grief Tech safety layers: strict provenance UI overlays (showing exact source texts for every generated answer) and clinical disclaimers to mitigate Prolonged Grief Disorder risks by explicitly framing the system as an archival tool.²
- **Phase 2: Secure Cloud MLOps & Collaboration (Months 4-9)** The system expands to support multi-user family collaboration. This requires deploying the Additively Homomorphic Encryption (AHE) pipeline, transitioning vector storage to a secure cloud database where similarity matching occurs on encrypted data.⁴ Distributed MLOps tools (MLflow) are integrated to handle model versioning, secure retraining workflows, and data lineage tracking. The RL-based adaptive governance controller is prototyped to manage privacy budgets.⁹
- **Phase 3: Quantum Research Sandbox (Months 10-15)** Dedicated solely to R&D, this phase establishes the Qiskit/Pennylane environment. Researchers begin offline benchmarking of amplitude encoding, mapping the classical 384-dimensional SBERT vectors into quantum states.⁷ The goal is to publish academic findings comparing quantum clustering efficiency and Parameterized Quantum Circuits (PQCs) against classical baseline metrics, establishing CHRONUS as a pioneer in applied QNLP.
- **Phase 4: Neuromorphic & Neuroscience Sandbox (Months 16-24)** Integration with Intel's Lava framework begins. Textual memories are mapped to Leaky Integrate-and-Fire (LIF) spike trains. The system simulates offline STDP "sleep" phases to study biologically plausible memory consolidation, evaluating how a simulated Hippocampal-Cortical (HiCo) network clusters related memories over time.⁶
- **Phase 5: Future Integration Layer (Years 2-5+)** Depending entirely on hardware maturity, successful algorithms from the research sandboxes are transitioned into the main product pipeline. This may involve shifting real-time semantic clustering to hybrid quantum-classical algorithms on quantum coprocessors, or integrating SNN associative recall mechanisms to dynamically re-weight memory retrieval based on biological Long-Term Potentiation (LTP) analogs.

9. Ethics and Governance

Because CHRONUS processes the intimate posthumous data of real individuals, ethical governance is the paramount design constraint.

- **Consent and Identity Rights:** The system must enforce cryptographic proof of consent before any persona generation begins, aligning with guidelines for pre-mortem legacy planning.¹³ The unauthorized third-party simulation of individuals is strictly prohibited at the application layer.
- **Data Provenance and Hallucination Control:** To preserve the integrity of the deceased, every generated response must display an unbroken data lineage to the original source document. The deterministic Mix Method eliminates ideological hallucination, preventing the AI from generating novel opinions that could distort the historical reality of the

individual.¹

- **Mitigation of Prolonged Grief Disorder (PGD):** The UX design must constantly counteract the illusion of sentience. By utilizing clinical UI markers, explicit citations, and session duration limits, the system actively prevents maladaptive attachments.² Escalation cues pointing to human mental health resources must be triggered if extreme emotional distress is detected in user queries.
- **Deletion and the Right to be Forgotten:** Users must possess the capability for immediate, unrecoverable deletion of all local vectors, trained models, and voice samples through a single application interface.
- **Privacy and Cloud Commodification:** The integration of the cloud tier strictly utilizes Additively Homomorphic Encryption (AHE) and Differential Privacy.⁵ This provides mathematical guarantees that no corporate entity, cloud provider, or malicious actor can reconstruct the plaintext data, read the intimate memories, or utilize the digital estate for foundational model training, effectively neutralizing the risk of commercial exploitation.²

10. Technical Stack Recommendations

To execute this architecture, the technology stack must be partitioned into Production and Research tiers, balancing immediate stability with future innovation.

What Stays (Production Core)

- **React 18 / TypeScript / Tailwind CSS:** Remains the optimal choice for a responsive, cross-platform user interface that handles complex workflows.¹
- **Python / Flask:** Continues to serve as the ideal glue layer for AI and cryptographic workloads.¹
- **Sentence-BERT (all-MiniLM-L6-v2):** Retained as the most computationally efficient dense embedding engine for edge and local deployments.¹
- **ChromaDB:** Retained strictly for the local-first execution tier due to its zero-dependency, rapid HNSW indexing.¹
- **XTTS-v2:** Retained as the primary local voice synthesizer, ensuring that biometric voice replication occurs without transmitting sensitive audio data over external APIs.¹

What Changes (Cloud / Scaling)

- **Cloud Vector Database:** Transition cloud deployments to a distributed vector store (e.g., Milvus or an AHE-enabled Postgres/pgvector instance). This database must be configured strictly to compute similarities on normalized encrypted data via PHE/AHE.⁴
- **Governance Engine:** Integrate Open Policy Agent (OPA) and Gatekeeper admission controllers to enforce strict access control, telemetry tracking, and zero-knowledge compliance verification in the cloud environment.²¹

What is Added (Research Sandboxes - Strictly Non-Production)

- **Lava Framework:** Intel's open-source library for programming neuromorphic architectures (Loihi 2). Used exclusively for offline, experimental memory consolidation

(STDP) research.¹⁸

- **Qiskit / PennyLane:** Integrated into a separate codebase to develop parameterized quantum circuits and simulate amplitude encoding for QNLP benchmarks.

11. Risks and Mitigations

Risk Domain	Identified Risk	Mitigation Strategy
Psychological	Maladaptive attachment / Prolonged Grief Disorder resulting from overly convincing illusions of sentience. ²	Enforce clinical UI design. Avoid generative "chit-chat" mechanisms. Keep all responses tethered strictly to historical text via the Mix Method and display visible source citations.
Security & Privacy	Catastrophic cloud vector database leakage exposing highly intimate personal data, PII, and credentials. ²⁰	Never transmit plaintext embeddings to the cloud. Implement Additively Homomorphic Encryption (AHE) to ensure cloud operations occur only on ciphertexts. ⁵
Technological (Quantum)	Current NISQ hardware noise destroying the fidelity of high-dimensional amplitude encoded embeddings. ⁷	Isolate QML entirely to an offline research sandbox. Rely strictly on classical SBERT for user-facing production systems until fault-tolerant hardware matures.
Technological (Neuromorphic)	Latency and hardware/software abstraction gaps preventing real-time integration of SNNs. ⁶	Utilize Lava/Loihi 2 solely for asynchronous backend "sleep phase" clustering and academic research, not for real-time conversational query answering. ⁶
Business / Operational	Massive computational overhead and latency	Shift cryptographic architecture from FHE to

	associated with Fully Homomorphic Encryption (FHE) in cloud deployments. ²³	PHE/AHE. By pre-normalizing vectors, the system only requires addition and scalar multiplication for dot products, drastically reducing latency. ⁴
--	--	---

12. Final Recommendation

To build CHRONUS into a system that satisfies academic rigor, product viability, and hackathon innovation, the strategy must embrace a **Tiered Reality Model**. Do not attempt to collapse the pillars into a singular, hype-driven narrative. If the goal is a hackathon win or a real startup, focus entirely on **Tier 1 (The Present Product)** and **Tier 2 (The Advanced Platform)**. Build the world’s most secure, ethically designed, local-first memory retrieval system. Use the Mix Method to eliminate hallucination and XTTS-v2 to guarantee vocal privacy. Scale this secure system to the cloud using Additively Homomorphic Encryption (AHE). This solves the distributed storage problem while providing a mathematical guarantee of privacy against data commodification, making the system immensely viable as a secure enterprise startup.²

Conversely, if the goal is an academic paper or a long-term research prototype, activate **Tier 3 (The Research Frontier)**. Treat Quantum AI and Neuromorphic Computing precisely as what they are: the frontier. Establish separate codebases (sandboxes) where classical data is mirrored and transformed. Use these sandboxes to study amplitude encoding ⁷ and biologically plausible memory semantization (HiCo SNN).⁶ This approach generates patents, produces high-impact papers, and prepares the CHRONUS architecture for the hardware convergence of the 2030s. By maintaining strict discipline between what runs in production today and what is modeled in research tomorrow, CHRONUS can uniquely dominate the digital legacy space without sacrificing scientific credibility.

13. Next-Step Implementation Checklist

Phase	Action Item	Owner / Role	Status
Immediate (Product)	Harden local ChromaDB pipeline and integrate explicit source-citation overlays into the	AI Systems Engineer / Frontend Dev	Pending

	React frontend.		
Immediate (Product)	Consult with clinical psychologists to draft and embed PGD-mitigation disclaimers into the UX. ²	Product Architect	Pending
Mid-Term (Cloud)	Prototype Additively Homomorphic Encryption (AHE) using Python libraries (e.g., LightPHE) to test encrypted dot products. ⁴	Cryptography / Systems Researcher	Pending
Mid-Term (Cloud)	Set up MLflow for local model versioning and data lineage tracking.	MLOps Engineer	Pending
Long-Term (Research)	Establish the Intel Lava framework environment and begin mapping SBERT outputs to LIF spike trains. ⁶	Neuromorphic Strategist	Pending
Long-Term (Research)	Write initial Qiskit scripts to test amplitude encoding of 16-dimensional reduced vectors on simulators. ⁷	Quantum Computing Researcher	Pending

Works cited

1. CHRONUS_Basic_Version_Report.pdf
2. Research Paper Generation and Refinement,
<https://drive.google.com/open?id=1U5Wz271HPAWP5S5us0l6pGiltVCrp4PHN1qo>

[pwnfUGw](#)

3. How AI Is Changing the Way Humans Grieve - School of Data Science, accessed July 13, 2026, <https://datascience.virginia.edu/news/how-ai-changing-way-humans-grieve>
4. Encrypted Vector Similarity Computations Using Partially Homomorphic Encryption: Applications and Performance Analysis - arXiv, accessed July 13, 2026, <https://arxiv.org/pdf/2503.05850>
5. arXiv:2502.14291v1 [cs.CR] 20 Feb 2025, accessed July 13, 2026, <https://arxiv.org/pdf/2502.14291>
6. Implementation of a Hippocampal-Cortical spiking neural network for memory semantization on Loihi 2 neuromorphic hardware - WebThesis, accessed July 13, 2026, <https://webthesis.biblio.polito.it/39912/1/tesi.pdf>
7. Quantum Natural Language Processing: A Comprehensive Survey of Models, Architectures, and Evaluation Methods - ACL Anthology, accessed July 13, 2026, <https://aclanthology.org/2025.quantumnlp-1.8.pdf>
8. Undergraduate Research - Fall 2025 Ronit Jha, Yanin Charoenpornasawat, Ayomide JeJe, Shane Sparks, Damir Cavar* Quantum AI and - AI & NLP Lab, accessed July 13, 2026, https://nlp-lab.org/publications/Quantum_Computing_For_AI_and_NLP.pdf
9. A Privacy-Preserving Cloud Architecture for Distributed Machine Learning at Scale - arXiv, accessed July 13, 2026, <https://arxiv.org/pdf/2512.10341>
10. Griefbots: Blurring the Reality of Death and the Illusion of Life - The University of Alabama at Birmingham, accessed July 13, 2026, <https://sites.uab.edu/humanrights/2025/02/07/griefbots-blurring-the-reality-of-death-and-the-illusion-of-life/>
11. Full article: The impact of continuing bonds following bereavement: A systematic review, accessed July 13, 2026, <https://www.tandfonline.com/doi/full/10.1080/07481187.2023.2223593>
12. Unfinished business, continuing bonds and bereaved college students' grief: a three-wave longitudinal study - PMC, accessed July 13, 2026, <https://pmc.ncbi.nlm.nih.gov/articles/PMC12239121/>
13. The Making of Digital Ghosts: Designing Ethical AI Afterlives - ResearchGate, accessed July 13, 2026, https://www.researchgate.net/publication/397984077_The_Making_of_Digital_Ghosts_Designing_Ethical_AI_Afterlives
14. The Making of Digital Ghosts: Designing Ethical AI Afterlives - arXiv, accessed July 13, 2026, <https://arxiv.org/pdf/2511.20094>
15. Quantum Natural Language Processing: A Comprehensive Review of Models, Methods, and Applications - arXiv, accessed July 13, 2026, <https://arxiv.org/html/2504.09909v2>
16. Quantum AI and NLP | springerprofessional.de, accessed July 13, 2026, <https://www.springerprofessional.de/quantum-ai-and-nlp/51904962>
17. Quantum Natural Language Processing - arXiv, accessed July 13, 2026, <https://arxiv.org/html/2403.19758v2>
18. Semantization of memories in a hippocampal-cortical spiking neural network |

Request PDF, accessed July 13, 2026,

https://www.researchgate.net/publication/391424672_Semantization_of_memories_in_a_hippocampal-cortical_spiking_neural_network

19. The generation and dynamic control of neural activity sequences - eDiss, accessed July 13, 2026, https://ediss.uni-goettingen.de/bitstream/handle/11858/15845/Lehr_PhD_Thesis.pdf?sequence=1&isAllowed=y
20. The AI Data You Forgot to Lock: How Exposed Vector Databases Put Organizations at Risk - Orca Security, accessed July 13, 2026, <https://orca.security/resources/blog/vector-database-security-risks/>
21. A Privacy-Preserving Cloud Architecture for Distributed Machine Learning at Scale - arXiv, accessed July 13, 2026, <https://arxiv.org/html/2512.10341v1>
22. A Privacy-Preserving Cloud Architecture for Distributed Machine Learning at Scale, accessed July 13, 2026, https://www.researchgate.net/publication/398602233_A_Privacy-Preserving_Cloud_Architecture_for_Distributed_Machine_Learning_at_Scale
23. A Note on Efficient Privacy-Preserving Similarity Search for Encrypted Vectors - arXiv, accessed July 13, 2026, <https://arxiv.org/html/2502.14291v1>
24. Encrypted Vector Similarity Computations Using Partially Homomorphic Encryption: Applications and Performance Analysis - arXiv, accessed July 13, 2026, <https://arxiv.org/html/2503.05850v1>
25. How does a distributed system for scalable vector databases work? - Reddit, accessed July 13, 2026, https://www.reddit.com/r/vectordatabase/comments/1jdikbs/how_does_a_distributed_system_for_scalable_vector/